

assignment-4

September 18, 2023

0.0.1 Dataset Inspection

```
[2]: import pandas as pd
import numpy as np
from sklearn import preprocessing
```

```
[3]: df = pd.read_csv("data.csv")
```

```
[19]: df.shape
```

```
[19]: (1599, 12)
```

```
[14]: df.head(20)
```

```
[14]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
0	7.4	0.700	0.00	1.9	0.076	
1	7.8	0.880	0.00	2.6	0.098	
2	7.8	0.760	0.04	2.3	0.092	
3	11.2	0.280	0.56	1.9	0.075	
4	7.4	0.700	0.00	1.9	0.076	
5	7.4	0.660	0.00	1.8	0.075	
6	7.9	0.600	0.06	1.6	0.069	
7	7.3	0.650	0.00	1.2	0.065	
8	7.8	0.580	0.02	2.0	0.073	
9	7.5	0.500	0.36	6.1	0.071	
10	6.7	0.580	0.08	1.8	0.097	
11	7.5	0.500	0.36	6.1	0.071	
12	5.6	0.615	0.00	1.6	0.089	
13	7.8	0.610	0.29	1.6	0.114	
14	8.9	0.620	0.18	3.8	0.176	
15	8.9	0.620	0.19	3.9	0.170	
16	8.5	0.280	0.56	1.8	0.092	
17	8.1	0.560	0.28	1.7	0.368	
18	7.4	0.590	0.08	4.4	0.086	
19	7.9	0.320	0.51	1.8	0.341	

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	\
0	11.0	34.0	0.9978	3.51	0.56	

1	25.0	67.0	0.9968	3.20	0.68
2	15.0	54.0	0.9970	3.26	0.65
3	17.0	60.0	0.9980	3.16	0.58
4	11.0	34.0	0.9978	3.51	0.56
5	13.0	40.0	0.9978	3.51	0.56
6	15.0	59.0	0.9964	3.30	0.46
7	15.0	21.0	0.9946	3.39	0.47
8	9.0	18.0	0.9968	3.36	0.57
9	17.0	102.0	0.9978	3.35	0.80
10	15.0	65.0	0.9959	3.28	0.54
11	17.0	102.0	0.9978	3.35	0.80
12	16.0	59.0	0.9943	3.58	0.52
13	9.0	29.0	0.9974	3.26	1.56
14	52.0	145.0	0.9986	3.16	0.88
15	51.0	148.0	0.9986	3.17	0.93
16	35.0	103.0	0.9969	3.30	0.75
17	16.0	56.0	0.9968	3.11	1.28
18	6.0	29.0	0.9974	3.38	0.50
19	17.0	56.0	0.9969	3.04	1.08

	alcohol	quality
0	9.4	5
1	9.8	5
2	9.8	5
3	9.8	6
4	9.4	5
5	9.4	5
6	9.4	5
7	10.0	7
8	9.5	7
9	10.5	5
10	9.2	5
11	10.5	5
12	9.9	5
13	9.1	5
14	9.2	5
15	9.2	5
16	10.5	7
17	9.3	5
18	9.0	4
19	9.2	6

```
[20]: df.describe()
```

```
[20]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar \
count	1599.000000	1599.000000	1599.000000	1599.000000
mean	8.319637	0.527821	0.270976	2.538806

std	1.741096	0.179060	0.194801	1.409928
min	4.600000	0.120000	0.000000	0.900000
25%	7.100000	0.390000	0.090000	1.900000
50%	7.900000	0.520000	0.260000	2.200000
75%	9.200000	0.640000	0.420000	2.600000
max	15.900000	1.580000	1.000000	15.500000

	chlorides	free sulfur dioxide	total sulfur dioxide	density \
count	1599.000000	1599.000000	1599.000000	1599.000000
mean	0.087467	15.874922	46.467792	0.996747
std	0.047065	10.460157	32.895324	0.001887
min	0.012000	1.000000	6.000000	0.990070
25%	0.070000	7.000000	22.000000	0.995600
50%	0.079000	14.000000	38.000000	0.996750
75%	0.090000	21.000000	62.000000	0.997835
max	0.611000	72.000000	289.000000	1.003690

	pH	sulphates	alcohol	quality
count	1599.000000	1599.000000	1599.000000	1599.000000
mean	3.311113	0.658149	10.422983	5.636023
std	0.154386	0.169507	1.065668	0.807569
min	2.740000	0.330000	8.400000	3.000000
25%	3.210000	0.550000	9.500000	5.000000
50%	3.310000	0.620000	10.200000	6.000000
75%	3.400000	0.730000	11.100000	6.000000
max	4.010000	2.000000	14.900000	8.000000

```
[6]: df.columns
```

```
[6]: Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',
          'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',
          'pH', 'sulphates', 'alcohol', 'quality'],
          dtype='object')
```

```
[11]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          1599 non-null   float64
1   volatile acidity       1599 non-null   float64
2   citric acid            1599 non-null   float64
3   residual sugar         1599 non-null   float64
4   chlorides              1599 non-null   float64
5   free sulfur dioxide    1599 non-null   float64
```

```

6  total sulfur dioxide  1599 non-null  float64
7  density              1599 non-null  float64
8  pH                  1599 non-null  float64
9  sulphates           1599 non-null  float64
10 alcohol             1599 non-null  float64
11 quality             1599 non-null  int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB

```

0.0.2 Data Pre-processing and Visualisation


```
[15]: import seaborn as sns
import matplotlib.pyplot as plt
```

```
[9]: df.isnull().values.any()
```

```
[9]: False
```

```
[26]: sns.barplot(x=df['quality'].value_counts().index, y=df['quality'].
↪value_counts())
```

```
[26]: <Axes: ylabel='quality'>
```



```
[28]: sns.distplot(df['alcohol'])
```

<ipython-input-28-570de8ff0310>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>


```
sns.distplot(df['alcohol'])
```

[28]: <Axes: xlabel='alcohol', ylabel='Density'>



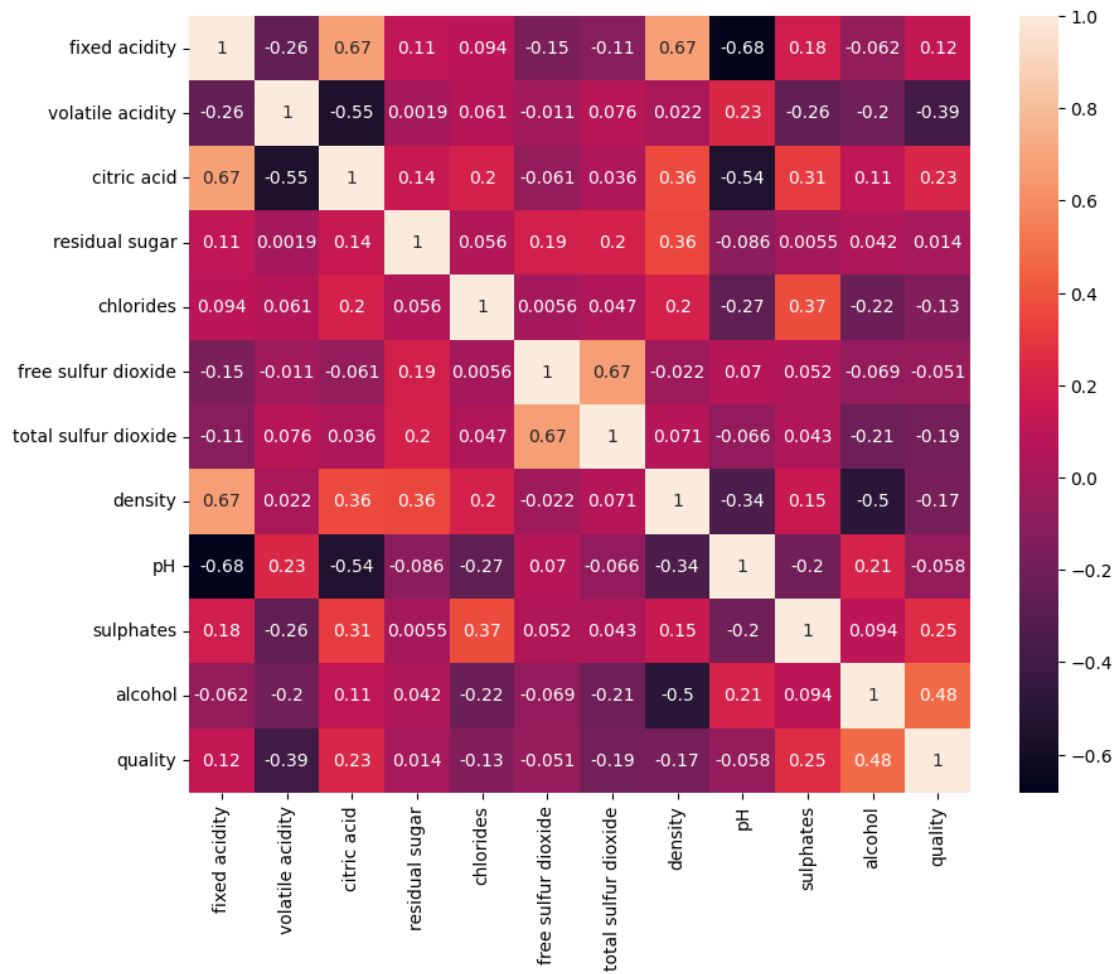
```
[29]: sns.lineplot(x=df['quality'],y=df['volatile acidity'])
```

[29]: <Axes: xlabel='quality', ylabel='volatile acidity'>



```
[31]: plt.figure(figsize=(10,8))  
      sns.heatmap(df.corr(), annot=True)
```

[31]: <Axes: >



0.0.3 Data Splitting

```
[34]: y = df['quality']
      y.head()
```

```
[34]: 0    5
      1    5
      2    5
      3    6
      4    5
      Name: quality, dtype: int64
```

```
[40]: x = df.drop(columns=['quality'], axis=1)
      x.head()
```

```
[40]: fixed acidity volatile acidity citric acid residual sugar chlorides \
0      7.4      0.70      0.00      1.9      0.076
1      7.8      0.88      0.00      2.6      0.098
2      7.8      0.76      0.04      2.3      0.092
3     11.2      0.28      0.56      1.9      0.075
4      7.4      0.70      0.00      1.9      0.076

      free sulfur dioxide total sulfur dioxide density pH sulphates \
0      11.0      34.0  0.9978 3.51      0.56
1      25.0      67.0  0.9968 3.20      0.68
2      15.0      54.0  0.9970 3.26      0.65
3      17.0      60.0  0.9980 3.16      0.58
4      11.0      34.0  0.9978 3.51      0.56

      alcohol
0      9.4
1      9.8
2      9.8
3      9.8
4      9.4
```

0.0.4 Scaling The Data

```
[41]: from sklearn.preprocessing import MinMaxScaler
```

```
[43]: x_scaled = pd.DataFrame(MinMaxScaler().fit_transform(x), columns=x.columns)
x_scaled.head()
```

```
[43]: fixed acidity volatile acidity citric acid residual sugar chlorides \
0      0.247788      0.397260      0.00      0.068493  0.106845
1      0.283186      0.520548      0.00      0.116438  0.143573
2      0.283186      0.438356      0.04      0.095890  0.133556
3      0.584071      0.109589      0.56      0.068493  0.105175
4      0.247788      0.397260      0.00      0.068493  0.106845

      free sulfur dioxide total sulfur dioxide density pH sulphates \
0      0.140845      0.098940  0.567548  0.606299  0.137725
1      0.338028      0.215548  0.494126  0.362205  0.209581
2      0.197183      0.169611  0.508811  0.409449  0.191617
3      0.225352      0.190813  0.582232  0.330709  0.149701
4      0.140845      0.098940  0.567548  0.606299  0.137725

      alcohol
0  0.153846
1  0.215385
2  0.215385
3  0.215385
```


4 0.153846

0.0.5 Train and Test Split

```
[44]: from sklearn.model_selection import train_test_split
```

```
[62]: x_train, x_test, y_train, y_test = train_test_split(x_scaled, y, test_size=0.4,
    random_state=0)
```

```
[63]: x_train.shape
```

[63]: (959, 11)

```
[64]: y_train.shape
```

```
[64]: (959,)
```

```
[65]: x_test.shape
```

[65]: (640, 11)

```
[66]: y_test.shape
```

[66]: (640,)

0.0.6 Logistic Model

```
[67]: from sklearn.linear_model import LogisticRegression
      model1 = LogisticRegression()
```

```
[68]: model1.fit(x_train,y_train)
```

[68]: LogisticRegression()

```
[69]: y_pred1 = model1.predict(x_test)
      y_pred1
```

```
[69]: array([6, 5, 6, 5, 6, 5, 5, 6, 5, 5, 5, 5, 6, 5, 6, 6, 7, 6, 6, 5, 6, 5,  
            6, 6, 5, 5, 5, 6, 5, 6, 6, 6, 6, 5, 6, 6, 5, 6, 6, 6, 5, 6, 7, 7,  
            6, 5, 5, 6, 6, 6, 5, 5, 6, 6, 6, 5, 5, 5, 7, 5, 5, 6, 6, 6, 5, 6,  
            5, 6, 6, 6, 5, 5, 5, 5, 5, 6, 5, 5, 5, 6, 6, 5, 6, 6, 5, 5, 6, 5,  
            5, 5, 5, 5, 6, 5, 6, 5, 6, 5, 5, 6, 7, 6, 6, 7, 6, 5, 6, 5, 6, 5,  
            6, 5, 6, 5, 6, 6, 6, 7, 6, 6, 5, 6, 5, 5, 6, 6, 5, 5, 6, 6, 5, 5,  
            6, 6, 6, 5, 6, 5, 6, 5, 6, 5, 5, 5, 6, 6, 6, 6, 6, 5, 6, 6, 5, 6,  
            6, 5, 5, 5, 6, 6, 6, 6, 6, 5, 6, 5, 6, 7, 5, 6, 6, 5, 5, 7, 6, 6,  
            6, 7, 6, 5, 5, 7, 5, 6, 7, 5, 5, 6, 5, 6, 6, 6, 5, 5, 5, 5, 5, 5])
```

```

5, 5, 5, 6, 5, 5, 5, 5, 5, 6, 6, 5, 6, 5, 5, 7, 5, 5, 6, 6, 6, 5,
5, 6, 6, 6, 5, 6, 7, 6, 5, 5, 5, 6, 5, 6, 6, 6, 6, 7, 7, 6, 5, 5,
5, 5, 6, 5, 5, 5, 5, 6, 5, 5, 5, 5, 5, 5, 5, 5, 6, 5, 7, 5, 5,
5, 5, 5, 5, 6, 7, 5, 6, 6, 6, 6, 6, 6, 5, 7, 6, 5, 7, 6, 6, 6, 5,
5, 5, 6, 6, 6, 6, 6, 5, 5, 6, 5, 5, 5, 5, 6, 5, 5, 5, 6, 6, 5, 5,
5, 5, 6, 6, 5, 5, 5, 7, 6, 6, 5, 7, 5, 5, 6, 6, 6, 5, 7, 5, 5, 5,
7, 5, 5, 6, 5, 6, 6, 5, 5, 5, 5, 5, 5, 6, 6, 5, 5, 5, 5, 6, 6, 6,
5, 6, 7, 5, 6, 6, 6, 6, 5, 6, 5, 6, 6, 5, 6, 6, 5, 5, 6, 6, 5, 6,
5, 5, 6, 6, 6, 6, 6, 6, 5, 5, 5, 6, 6, 5, 6, 5, 7, 5, 5, 7, 5, 6,
5, 6, 6, 5, 5, 5, 5, 6, 5, 6, 6, 5, 6, 5, 5, 5, 6, 6, 5, 6, 5,
5, 6, 6, 6, 7, 6, 5, 6, 6, 6, 5, 5, 5, 6, 5, 6, 6, 7, 6, 6, 5, 5,
6, 5, 6, 5, 6, 5, 5, 7, 5, 6, 6, 5, 6, 7, 6, 5, 6, 5, 6, 5, 5, 7,
5, 5, 5, 5, 6, 5, 5, 5, 6, 5, 5, 5, 6, 5, 5, 5, 6, 5, 6, 7, 5, 6,
6, 7, 6, 5, 6, 5, 5, 6, 5, 6, 6, 6, 5, 5, 6, 5, 6, 5, 5, 5, 5, 6,
5, 5, 6, 6, 6, 6, 5, 5, 7, 6, 6, 5, 5, 5, 5, 5, 6, 5, 5, 5, 6, 5,
6, 6, 6, 5, 6, 7, 5, 6, 6, 6, 6, 6, 5, 5, 6, 6, 6, 7, 6, 7, 5,
6, 6, 5, 6, 6, 5, 6, 7, 6, 5, 6, 6, 5, 5, 5, 6, 6, 7, 6, 5, 6, 5,
5, 6, 6, 6, 5, 6, 6, 6, 6, 6, 6, 5, 5, 5, 5, 6, 7, 6, 5, 6, 6,
6, 6, 6, 6, 5, 6, 6, 5, 7, 6, 6, 5, 6, 6, 5, 6, 5, 5, 6, 5, 6, 6,
6, 5, 5, 5, 6, 5, 6, 6, 6, 6, 5, 5, 5, 5, 6, 5, 6, 6, 5, 6, 5, 5,
6, 6]

```

0.0.7 Evaluation

```
[70]: from sklearn.metrics import accuracy_score, classification_report
```

```
[71]: acc1 = accuracy_score(y_test, y_pred1)
acc1
```

```
[71]: 0.615625
```

```
[73]: pd.crosstab(y_test, y_pred1)
```

```
[73]: col_0      5      6      7
quality
3          4      0      0
4         11     12      0
5        208     70      1
6         76    170     22
7           3     41     16
8           0      3      3
```

```
[74]: print(classification_report(y_test, y_pred1))
```

```

              precision    recall  f1-score   support

3             0.00         0.00         0.00         4

```

4	0.00	0.00	0.00	23
5	0.69	0.75	0.72	279
6	0.57	0.63	0.60	268
7	0.38	0.27	0.31	60
8	0.00	0.00	0.00	6
accuracy			0.62	640
macro avg	0.27	0.27	0.27	640
weighted avg	0.58	0.62	0.59	640

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```