

assignment-5

October 5, 2023

1 Kaggle Connection & DataFrame setup

```
[601]: !pip install -q kaggle
```

```
[602]: !mkdir ~/.kaggle
```

mkdir: cannot create directory '/root/.kaggle': File exists

```
[603]: !cp kaggle.json ~/.kaggle
```

```
[604]: ! kaggle datasets download -d vjchoudhary7/  
↪customer-segmentation-tutorial-in-python
```

Warning: Your Kaggle API key is readable by other users on this system! To fix this, you can run 'chmod 600 /root/.kaggle/kaggle.json'

Downloading customer-segmentation-tutorial-in-python.zip to /content

0% 0.00/1.55k [00:00<?, ?B/s]

100% 1.55k/1.55k [00:00<00:00, 2.51MB/s]

```
[605]: !unzip /content/customer-segmentation-tutorial-in-python.zip
```

Archive: /content/customer-segmentation-tutorial-in-python.zip
inflating: Mall_Customers.csv

2 Pre-Processing

```
[606]: import pandas as pd  
import numpy as np  
import seaborn as sns  
import matplotlib.pyplot as plt
```

```
[607]: df = pd.read_csv('./Mall_Customers.csv')  
df.head()
```

```
[607]:
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	39
1	2	Male	21	15	81

2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40

```
[608]: df.describe()
```

```
[608]:
```

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)
count	200.000000	200.000000	200.000000	200.000000
mean	100.500000	38.850000	60.560000	50.200000
std	57.879185	13.969007	26.264721	25.823522
min	1.000000	18.000000	15.000000	1.000000
25%	50.750000	28.750000	41.500000	34.750000
50%	100.500000	36.000000	61.500000	50.000000
75%	150.250000	49.000000	78.000000	73.000000
max	200.000000	70.000000	137.000000	99.000000

```
[609]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CustomerID            200 non-null   int64
1   Gender                 200 non-null   object
2   Age                    200 non-null   int64
3   Annual Income (k$)     200 non-null   int64
4   Spending Score (1-100) 200 non-null   int64
dtypes: int64(4), object(1)
memory usage: 7.9+ KB
```

```
[610]: df.isnull().values.any()
```

```
[610]: False
```

```
[611]: df.shape
```

```
[611]: (200, 5)
```


```
[612]: # Dropping 'CustomerID' as it has no impact or connection to dataset or data_
      ↪ values
df.drop(['CustomerID'], axis=1, inplace=True)
```

```
[613]: sns.heatmap(df.corr(), annot=True)
```

```
<ipython-input-613-6dc1c4c1753e>:1: FutureWarning: The default value of
numeric_only in DataFrame.corr is deprecated. In a future version, it will
default to False. Select only valid columns or specify the value of numeric_only
```


```
to silence this warning.  
sns.heatmap(df.corr(), annot=True)
```

[613]: <Axes: >



```
[614]: sns.pairplot(df)
```

[614]: <seaborn.axisgrid.PairGrid at 0x7daddca476a0>



3 Converting Categorical Data (Columns) to Numerical

```
[615]: df['Gender'].value_counts()
```

```
[615]: Female    112
      Male      88
      Name: Gender, dtype: int64
```

```
[616]: from sklearn.preprocessing import LabelEncoder
      le = LabelEncoder()
```

```
[617]: # Label Encoding 'Gender' column
# '1' == 'Male' && 0 == 'Female'
df['Gender'] = le.fit_transform(df.Gender)
```

```
[618]: df.head()
```

```
[618]:   Gender  Age  Annual Income (k$)  Spending Score (1-100)
0       1   19                15                39
1       1   21                15                81
2       0   20                16                 6
3       0   23                16               77
4       0   31                17               40
```


```
[619]: df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Gender                200 non-null   int64
1   Age                   200 non-null   int64
2   Annual Income (k$)    200 non-null   int64
3   Spending Score (1-100) 200 non-null   int64
dtypes: int64(4)
memory usage: 6.4 KB
```


4 Data Analysis, Outlier Detection & Outlier Elimination

```
[620]: sns.displot(df['Age'])
sns.displot(df['Annual Income (k$)'])
sns.displot(df['Spending Score (1-100)'])
```

```
[620]: <seaborn.axisgrid.FacetGrid at 0x7daddba3dae0>
```







```
[621]: sns.distplot(df['Age'])
```

<ipython-input-621-0fafa04ea3f6>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['Age'])
```

```
[621]: <Axes: xlabel='Age', ylabel='Density'>
```

```
[622]: sns.distplot(df['Annual Income (k$)'])
```

<ipython-input-622-5c9bfeb4bab1>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['Annual Income (k$)'])
```

```
[622]: <Axes: xlabel='Annual Income (k$)', ylabel='Density'>
```



```
[623]: sns.distplot(df['Spending Score (1-100)'])
```

<ipython-input-623-beed7b40d5ab>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>


```
sns.distplot(df['Spending Score (1-100)'])
```

```
[623]: <Axes: xlabel='Spending Score (1-100)', ylabel='Density'>
```



```
[624]: sns.boxplot(df['Annual Income (k$)'])
```

```
[624]: <Axes: >
```



```
[625]: Q1 = df['Annual Income (k$)'].quantile(0.25)
       Q3 = df['Annual Income (k$)'].quantile(0.75)
```


```
[626]: IQR = Q3 - Q1
       whisker_width = 1.5
```

```
[627]: lower_whisker = Q1 - (whisker_width*IQR)
       upper_whisker = Q3 + (whisker_width*IQR)
```

```
[628]: df['Annual Income (k$)'] = np.where(df['Annual Income (k$)'] > upper_whisker,
      ↪upper_whisker, np.where(df['Annual Income (k$)'] < lower_whisker,
      ↪lower_whisker, df['Annual Income (k$)']))
```


```
[629]: sns.boxplot(df['Annual Income (k$)'])
```

```
[629]: <Axes: >
```



```
[630]: plt.scatter(df['Annual Income (k$)'], df['Spending Score (1-100)'])
```

```
[630]: <matplotlib.collections.PathCollection at 0x7daddb65d7b0>
```



```
[631]: from sklearn.preprocessing import MinMaxScaler
```

```
[632]: X_train = df.drop(['Spending Score (1-100)'], axis=1)
X_train = pd.DataFrame(MinMaxScaler().fit_transform(X_train), columns=X_train.
    ↳ columns)

Y_train = df['Spending Score (1-100)']
```

```
[633]: X_train.head(), Y_train.head()
```

```
[633]: (   Gender      Age  Annual Income (k$)
0      1.0  0.019231          0.000000
1      1.0  0.057692          0.000000
2      0.0  0.038462          0.008493
3      0.0  0.096154          0.008493
4      0.0  0.250000          0.016985,
0      39
1      81
2       6
3      77
4      40
Name: Spending Score (1-100), dtype: int64)
```

5 Finding Elbow Point (Possible 'K' value)

```
[634]: from sklearn.cluster import KMeans
```

```
[635]: k_rng = range(1,40)
sse = []

for k in k_rng:
    km = KMeans(n_clusters=k)
    km.fit(df[['Annual Income (k$)', 'Spending Score (1-100)']])
    sse.append(km.inertia )
```

```

/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n init` explicitly to suppress the warning
    warnings.warn(

```



```

warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(

```


[636]: sse

[636]: [268717.56374999999,
180456.95393668828,
105529.63651521516,
72861.0524925923,
43639.878158556414,
36620.76718928144,
29829.066471160957,
24564.382281135964,
21382.608462322565,
19244.814450778555,
17205.548702513755,
15502.186514191457,
13958.923407658516,
12628.419801889288,
11912.3291765694,
10619.603376831503,
9803.323762719561,
8965.422420223756,
8637.61374868696,

```
7904.32068998865,  
7385.5635939800795,  
7173.877767501928,  
6543.25370736201,  
6410.3651709766855,  
5833.792235750361,  
5613.621778221779,  
5360.4049035807175,  
5023.465877329192,  
4887.764547258296,  
4447.030006105006,  
4149.634623015872,  
4064.836663336663,  
3972.645093795093,  
3663.232523726274,  
3537.3413149350654,  
3272.340124458874,  
3107.4417117604617,  
3070.0163149350656,  
2991.7733901515153]
```

```
[637]: plt.xlabel('K')  
plt.ylabel('Sum of Squared Error')  
plt.plot(k_rng, sse)
```

```
[637]: [<matplotlib.lines.Line2D at 0x7daddb6e1540>]
```



```
[638]: kmeans = KMeans(n_clusters=5)
kmeans
```

```
[638]: KMeans(n_clusters=5)
```

```
[639]: z = kmeans.fit_predict(df[['Annual Income (k$)', 'Spending Score (1-100)']])
      z
```

```

/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(

```

```
[639]: array([3, 1, 3, 1, 3, 1, 3, 1, 3, 1, 3, 1, 3, 1, 3, 1, 3, 1, 3, 1, 3, 1,
        3, 1, 3, 1, 3, 1, 3, 1, 3, 1, 3, 1, 3, 1, 3, 1, 3, 1, 3, 4,
        3, 1, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4,
        4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4,
        4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4,
        4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4,
        4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 2, 0, 2, 4, 2, 0, 2, 0, 2,
        4, 2, 0, 2, 0, 2, 0, 2, 0, 2, 4, 2, 0, 2, 0, 2, 0, 2, 0, 2, 0, 2,
        0, 2, 0, 2, 0, 2, 0, 2, 0, 2, 0, 2, 0, 2, 0, 2, 0, 2, 0, 2,
        0, 2, 0, 2, 0, 2, 0, 2, 0, 2, 0, 2, 0, 2, 0, 2, 0, 2, 0, 2,
```

```
0, 2], dtype=int32)
```

```
[640]: df_2 = df
```

```
[641]: df_2['Cluster'] = z
df_2.head()
```

```
[641]:
```

	Gender	Age	Annual Income (k\$)	Spending Score (1-100)	Cluster
0	1	19	15.0	39	3
1	1	21	15.0	81	1
2	0	20	16.0	6	3
3	0	23	16.0	77	1
4	0	31	17.0	40	3

```
[642]: plt.figure(figsize=(10,8))


df1 = df_2[df_2.Cluster==0]
df2 = df_2[df_2.Cluster==1]
df3 = df_2[df_2.Cluster==2]
df4 = df_2[df_2.Cluster==3]
df5 = df_2[df_2.Cluster==4]

plt.scatter(df1['Annual Income (k$)'], df1['Spending Score (1-100)'],
            ↪label='Cluster 1')
plt.scatter(df2['Annual Income (k$)'], df2['Spending Score (1-100)'],
            ↪label='Cluster 2')
plt.scatter(df3['Annual Income (k$)'], df3['Spending Score (1-100)'],
            ↪label='Cluster 3')
plt.scatter(df4['Annual Income (k$)'], df4['Spending Score (1-100)'],
            ↪label='Cluster 4')
plt.scatter(df5['Annual Income (k$)'], df5['Spending Score (1-100)'],
            ↪label='Cluster 5')

plt.scatter(kmeans.cluster_centers_[ :,0], kmeans.cluster_centers_[ :,1],
            ↪color='black', marker='*', label='Centroid')

plt.xlabel('Annual Income (k$)')
plt.ylabel('Spending Score (1-100)')
plt.legend()
```

```
[642]: <matplotlib.legend.Legend at 0x7daddb701b10>
```



6 Scaling & Train - Test Split

```
[643]: from sklearn.model_selection import train_test_split
```

```
[664]: X = df.drop(['Spending Score (1-100)'], axis=1)
X = pd.DataFrame(MinMaxScaler().fit_transform(X), columns=X.columns)
X.drop(['Cluster', 'Gender'], axis=1, inplace=True)
```

```
[665]: y = df['Spending Score (1-100)']
```

```
[666]: X.head(), y.head()
```

```
[666]: (
  Age  Annual Income (k$)
0  0.019231             0.000000
1  0.057692             0.000000
2  0.038462             0.008493
3  0.096154             0.008493
4  0.250000             0.016985,
```

```
0    39
1    81
2     6
3    77
4    40
Name: Spending Score (1-100), dtype: int64)
```

7 KNN and Logistic Regression Modeling

```
[667]: from sklearn.linear_model import LinearRegression, LogisticRegression
       from sklearn.neighbors import KNeighborsRegressor
```

```
[668]: lr = LogisticRegression(max_iter=10000)
       knn = KNeighborsRegressor(n_neighbors=3)
```

7.1 KNN Model

```
[669]: xtrain, xtest, ytrain, ytest = train_test_split(X, y, test_size=0.25,
       ↪random_state=100)
```

```
[670]: xtest.shape, xtrain.shape
```

```
[670]: ((50, 2), (150, 2))
```

```
[671]: ytest.shape, ytrain.shape
```

```
[671]: ((50,), (150,))
```

```
[672]: knn.fit(xtrain, ytrain)
```

```
[672]: KNeighborsRegressor(n_neighbors=3)
```

```
[673]: acc = knn.score(xtest, ytest)
       print(f"Accuracy for the KNN model is {acc*100:.2f}%")
```

Accuracy for the KNN model is 56.86%

7.2 Logistic Regression Model

```
[674]: Xtrain, Xtest, Ytrain, Ytest = train_test_split(X, y, test_size=0.30,
       ↪random_state=42)
```

```
[675]: Xtest.shape, Xtrain.shape
```

```
[675]: ((60, 2), (140, 2))
```

```
[676]: Ytest.shape, Ytrain.shape
```

```
[676]: ((60,), (140,))
```

```
[677]: lr.fit(Xtrain, Ytrain)
```

```
[677]: LogisticRegression(max_iter=10000)
```

```
[678]: acc = lr.score(Xtest, Ytest)
print(f"Accuracy for the Logistic Regression model is {acc*100:.2f}%")
```

Accuracy for the Logistic Regression model is 0.00%

8 Prediction

```
[691]: prediction1 = knn.predict([[36 , 24]])[0]
prediction2 = lr.predict([[36 , 24]])[0]
```

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but KNeighborsRegressor was fitted with feature names

```
warnings.warn(
```

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names

```
warnings.warn(
```

```
[696]: print("KNN Model")
print(f"Age: 36, Salary(k$): 24.0, Spending Score(1-100): {prediction1:.2f}")

print('\n', "*" * 10, '\n')

print("Logistic Regression Model")
print(f"Age: 36, Salary(k$): 24.0, Spending Score(1-100): {prediction2:.2f}")
```

KNN Model

Age: 36, Salary(k\$): 24.0, Spending Score(1-100): 37.33

Logistic Regression Model

Age: 36, Salary(k\$): 24.0, Spending Score(1-100): 43.00