

assignment-5

October 3, 2023

1 Kaggle Connection & DataFrame setup

```
[451]: !pip install -q kaggle
```

```
[452]: !mkdir ~/.kaggle
```

mkdir: cannot create directory '/root/.kaggle': File exists

```
[453]: !cp kaggle.json ~/.kaggle
```

```
[454]: ! kaggle datasets download -d vjchoudhary7/  
↪customer-segmentation-tutorial-in-python
```

Warning: Your Kaggle API key is readable by other users on this system! To fix this, you can run 'chmod 600 /root/.kaggle/kaggle.json'

Downloading customer-segmentation-tutorial-in-python.zip to /content

0% 0.00/1.55k [00:00<?, ?B/s]

100% 1.55k/1.55k [00:00<00:00, 4.28MB/s]

```
[455]: !unzip /content/customer-segmentation-tutorial-in-python.zip
```

Archive: /content/customer-segmentation-tutorial-in-python.zip
inflating: Mall_Customers.csv

2 Pre-Processing

```
[456]: import pandas as pd  
import numpy as np  
import seaborn as sns  
import matplotlib.pyplot as plt
```

```
[457]: df = pd.read_csv('./Mall_Customers.csv')  
df.head()
```

```
[457]:
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	39
1	2	Male	21	15	81

2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40

```
[458]: df.describe()
```

```
[458]:
```

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)
count	200.000000	200.000000	200.000000	200.000000
mean	100.500000	38.850000	60.560000	50.200000
std	57.879185	13.969007	26.264721	25.823522
min	1.000000	18.000000	15.000000	1.000000
25%	50.750000	28.750000	41.500000	34.750000
50%	100.500000	36.000000	61.500000	50.000000
75%	150.250000	49.000000	78.000000	73.000000
max	200.000000	70.000000	137.000000	99.000000

```
[459]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CustomerID            200 non-null   int64
1   Gender                200 non-null   object
2   Age                   200 non-null   int64
3   Annual Income (k$)    200 non-null   int64
4   Spending Score (1-100) 200 non-null   int64
dtypes: int64(4), object(1)
memory usage: 7.9+ KB
```

```
[460]: df.isnull().values.any()
```

```
[460]: False
```

```
[461]: df.shape
```

```
[461]: (200, 5)
```


```
[462]: # Dropping 'CustomerID' as it has no impact or connection to dataset or data_
      ↪ values
df.drop(['CustomerID'], axis=1, inplace=True)
```

```
[463]: sns.heatmap(df.corr(), annot=True)
```

```
<ipython-input-463-6dc1c4c1753e>:1: FutureWarning: The default value of
numeric_only in DataFrame.corr is deprecated. In a future version, it will
default to False. Select only valid columns or specify the value of numeric_only
```


```
to silence this warning.  
sns.heatmap(df.corr(), annot=True)
```

[463]: <Axes: >



```
[464]: sns.pairplot(df)
```

[464]: <seaborn.axisgrid.PairGrid at 0x7b17495fb400>



3 Converting Categorical Data (Columns) to Numerical

```
[465]: df['Gender'].value_counts()
```

```
[465]: Female    112
      Male      88
      Name: Gender, dtype: int64
```

```
[466]: from sklearn.preprocessing import LabelEncoder
      le = LabelEncoder()
```

```
[467]: # Label Encoding 'Gender' column
# '1' == 'Male' && 0 == 'Female'
df['Gender'] = le.fit_transform(df.Gender)
```

```
[468]: df.head()
```

```
[468]:   Gender  Age  Annual Income (k$)  Spending Score (1-100)
0       1   19                15                39
1       1   21                15                81
2       0   20                16                 6
3       0   23                16               77
4       0   31                17               40
```


```
[469]: df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Gender                200 non-null   int64
1   Age                   200 non-null   int64
2   Annual Income (k$)    200 non-null   int64
3   Spending Score (1-100) 200 non-null   int64
dtypes: int64(4)
memory usage: 6.4 KB
```


4 Data Analysis, Outlier Detection & Outlier Elimination

```
[470]: sns.displot(df['Age'])
sns.displot(df['Annual Income (k$)'])
sns.displot(df['Spending Score (1-100)'])
```

```
[470]: <seaborn.axisgrid.FacetGrid at 0x7b1748477ca0>
```







```
[471]: sns.distplot(df['Age'])
```

<ipython-input-471-0fafe04ea3f6>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['Age'])
```

```
[471]: <Axes: xlabel='Age', ylabel='Density'>
```

```
[472]: sns.distplot(df['Annual Income (k$)'])
```

<ipython-input-472-5c9bfeb4bab1>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['Annual Income (k$)'])
```

```
[472]: <Axes: xlabel='Annual Income (k$)', ylabel='Density'>
```



```
[473]: sns.distplot(df['Spending Score (1-100)'])
```

<ipython-input-473-beed7b40d5ab>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>


```
sns.distplot(df['Spending Score (1-100)'])
```

```
[473]: <Axes: xlabel='Spending Score (1-100)', ylabel='Density'>
```



```
[474]: sns.boxplot(df['Annual Income (k$)'])
```

```
[474]: <Axes: >
```



```
[475]: Q1 = df['Annual Income (k$)'].quantile(0.25)
      Q3 = df['Annual Income (k$)'].quantile(0.75)
```


```
[476]: IQR = Q3 - Q1
      whisker_width = 1.5
```

```
[477]: lower_whisker = Q1 - (whisker_width*IQR)
      upper_whisker = Q3 + (whisker_width*IQR)
```

```
[478]: df['Annual Income (k$)'] = np.where(df['Annual Income (k$)'] > upper_whisker,
      ↪upper_whisker, np.where(df['Annual Income (k$)'] < lower_whisker,
      ↪lower_whisker, df['Annual Income (k$)']))
```


```
[479]: sns.boxplot(df['Annual Income (k$)'])
```

```
[479]: <Axes: >
```



```
[480]: plt.scatter(df['Annual Income (k$)'], df['Spending Score (1-100)'])
```

```
[480]: <matplotlib.collections.PathCollection at 0x7b174809ea40>
```



```
[481]: X_train = df.drop(['Spending Score (1-100)'], axis=1)
       Y_train = df['Spending Score (1-100)']
```

```
[482]: X_train.head(), Y_train.head()
```

```
[482]: (   Gender  Age  Annual Income (k$)
0        1   19             15.0
1        1   21             15.0
2        0   20             16.0
3        0   23             16.0
4        0   31             17.0,
0       39
1       81
2        6
3       77
4       40
Name: Spending Score (1-100), dtype: int64)
```

5 Finding Elbow Point (Possible 'K' value)

```
[483]: from sklearn.cluster import KMeans
```


```
[ ]: k_rng = range(1,40)
     sse = []

     for k in k_rng:
         km = KMeans(n_clusters=k)
         km.fit(df[['Annual Income (k$)', 'Spending Score (1-100)']])
         sse.append(km.inertia_)
```

```
[ ]: sse
```

```
[486]: plt.xlabel('K')
     plt.ylabel('Sum of Squared Error')
     plt.plot(k_rng, sse)
```

```
[486]: [<matplotlib.lines.Line2D at 0x7b17481235b0>]
```




```


plt.scatter(df2['Annual Income (k$)'], df2['Spending Score (1-100)'],
            ↪label='Cluster 2')
plt.scatter(df3['Annual Income (k$)'], df3['Spending Score (1-100)'],
            ↪label='Cluster 3')
plt.scatter(df4['Annual Income (k$)'], df4['Spending Score (1-100)'],
            ↪label='Cluster 4')
plt.scatter(df5['Annual Income (k$)'], df5['Spending Score (1-100)'],
            ↪label='Cluster 5')

plt.scatter(kmeans.cluster_centers_[0], kmeans.cluster_centers_[1],
            ↪color='black', marker='*', label='Centroid')

plt.xlabel('Annual Income (k$)')
plt.ylabel('Spending Score (1-100)')
plt.legend()

```

[491]: <matplotlib.legend.Legend at 0x7b1747fbfd90>



6 Train - Test Split

```
[492]: df.drop(['Cluster'], axis=1, inplace=True)
```

```
[493]: df.head()
```

```
[493]:
```

	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	19	15.0	39
1	1	21	15.0	81
2	0	20	16.0	6
3	0	23	16.0	77
4	0	31	17.0	40

```
[494]: from sklearn.model_selection import train_test_split
```

```
[495]: X = df.drop(['Spending Score (1-100)'], axis=1)
y = df['Spending Score (1-100)']
```

7 KNN and Logistic Regression Modeling

```
[496]: from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
```

7.1 KNN Model

```
[554]: lr = LogisticRegression(max_iter=10000)
knn = KNeighborsClassifier(n_neighbors=3)
```

```
[555]: xtrain, xtest, ytrain, ytest = train_test_split(X, y, test_size=0.3,
↳ random_state=12)
```

```
[556]: xtest.shape, xtrain.shape
```

```
[556]: ((60, 3), (140, 3))
```

```
[557]: ytest.shape, ytrain.shape
```

```
[557]: ((60,), (140,))
```

```
[558]: knn.fit(xtrain, ytrain)
```

```
[558]: KNeighborsClassifier(n_neighbors=3)
```

```
[559]: acc = knn.score(xtest, ytest)
print(f"Accuracy for the KNN model is {acc*100:.2f}%")
```

Accuracy for the KNN model is 3.33%

7.2 Logistic Regression Model

```
[560]: Xtrain, Xtest, Ytrain, Ytest = train_test_split(X, y, test_size=0.3,  
↳ random_state=10)
```

```
[561]: Xtest.shape, Xtrain.shape
```

```
[561]: ((60, 3), (140, 3))
```

```
[562]: Ytest.shape, Ytrain.shape
```

```
[562]: ((60,), (140,))
```

```
[563]: lr.fit(Xtrain, Ytrain)
```

```
[563]: LogisticRegression(max_iter=10000)
```

```
[564]: acc = lr.score(xtest, ytest)  
print(f"Accuracy for the Logistic Regression model is {acc*100:.2f}%")
```

Accuracy for the Logistic Regression model is 18.33%

8 Prediction

```
[565]: prediction1 = knn.predict([[1.0, 36.0 , 24.0]])[0]  
prediction2 = lr.predict([[1.0, 36.0 , 24.0]])[0]
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does  
not have valid feature names, but KNeighborsClassifier was fitted with feature  
names
```

```
warnings.warn(  

```

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does  
not have valid feature names, but LogisticRegression was fitted with feature  
names
```

```
warnings.warn(  

```

```
[566]: print("KNN Model")  
print("Gender: Male, Age: 36, Salary(k$): 24.0, Spending Score(1-100): {}".  
↳ format(prediction1))  
  
print("Logistic Regression Model")  
print("Gender: Male, Age: 36, Salary(k$): 24.0, Spending Score(1-100): {}".  
↳ format(prediction2))
```

KNN Model

Gender: Male, Age: 36, Salary(k\$): 24.0, Spending Score(1-100): 35

Logistic Regression Model

Gender: Male, Age: 36, Salary(k\$): 24.0, Spending Score(1-100): 73

[]: