

jce5amulc

October 1, 2023

1 Kaggle Connection & DataFrame setup

```
[380]: !pip install -q kaggle
```

```
[381]: !mkdir ~/.kaggle
```

mkdir: cannot create directory '/root/.kaggle': File exists

```
[382]: !cp kaggle.json ~/.kaggle
```

```
[383]: ! kaggle datasets download -d vjchoudhary7/  
↪customer-segmentation-tutorial-in-python
```

Warning: Your Kaggle API key is readable by other users on this system! To fix this, you can run 'chmod 600 /root/.kaggle/kaggle.json'
customer-segmentation-tutorial-in-python.zip: Skipping, found more recently modified local copy (use --force to force download)

```
[384]: !unzip /content/customer-segmentation-tutorial-in-python.zip
```

Archive: /content/customer-segmentation-tutorial-in-python.zip
replace Mall_Customers.csv? [y]es, [n]o, [A]ll, [N]one, [r]ename: n

2 Pre-Processing

```
[385]: import pandas as pd  
import numpy as np  
import seaborn as sns  
import matplotlib.pyplot as plt
```

```
[386]: df = pd.read_csv('./Mall_Customers.csv')  
df.head()
```

```
[386]:
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6

3	4	Female	23	16	77
4	5	Female	31	17	40

```
[387]: df.describe()
```

```
[387]:
```

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)
count	200.000000	200.000000	200.000000	200.000000
mean	100.500000	38.850000	60.560000	50.200000
std	57.879185	13.969007	26.264721	25.823522
min	1.000000	18.000000	15.000000	1.000000
25%	50.750000	28.750000	41.500000	34.750000
50%	100.500000	36.000000	61.500000	50.000000
75%	150.250000	49.000000	78.000000	73.000000
max	200.000000	70.000000	137.000000	99.000000

```
[388]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CustomerID            200 non-null   int64
1   Gender                 200 non-null   object
2   Age                    200 non-null   int64
3   Annual Income (k$)     200 non-null   int64
4   Spending Score (1-100) 200 non-null   int64
dtypes: int64(4), object(1)
memory usage: 7.9+ KB
```

```
[389]: df.isnull().values.any()
```

```
[389]: False
```

```
[390]: df.shape
```

```
[390]: (200, 5)
```


```
[391]: # Dropping 'CustomerID' as it has no impact or connection to dataset or data
      ↪ values
      df.drop(['CustomerID'], axis=1, inplace=True)
```

```
[392]: sns.heatmap(df.corr(), annot=True)
```

```
<ipython-input-392-6dc1c4c1753e>:1: FutureWarning: The default value of
numeric_only in DataFrame.corr is deprecated. In a future version, it will
default to False. Select only valid columns or specify the value of numeric_only
```

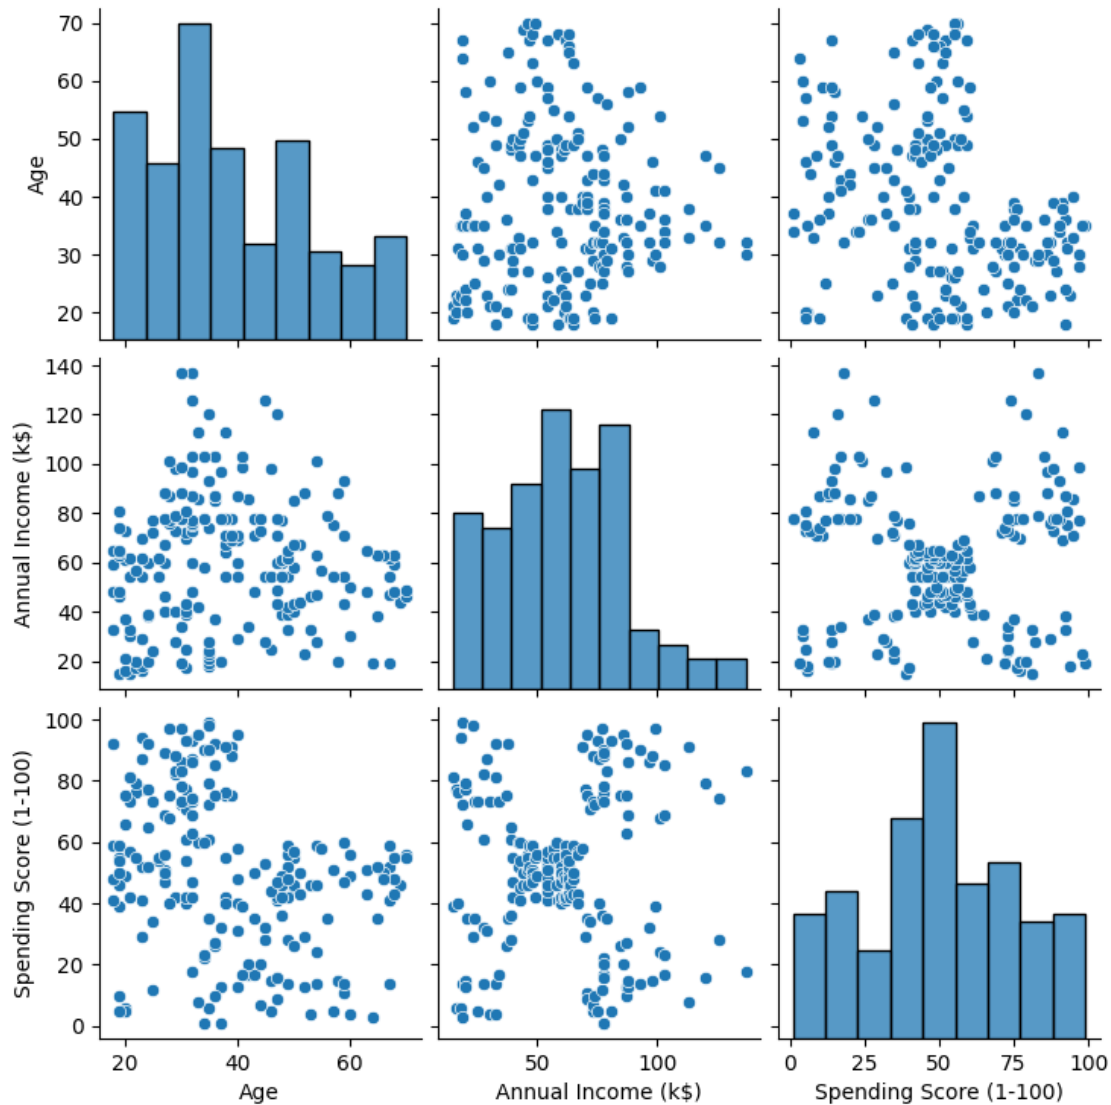
```
to silence this warning.  
sns.heatmap(df.corr(), annot=True)
```

[392]: <Axes: >



```
[393]: sns.pairplot(df)
```

[393]: <seaborn.axisgrid.PairGrid at 0x7c71dd8216f0>



3 Converting Categorical Data (Columns) to Numerical

```
[394]: df['Gender'].value_counts()
```

```
[394]: Female    112
      Male      88
      Name: Gender, dtype: int64
```

```
[395]: from sklearn.preprocessing import LabelEncoder
      le = LabelEncoder()
```

```
[396]: # Label Encoding 'Gender' column
# '1' == 'Male' && 0 == 'Female'
df['Gender'] = le.fit_transform(df.Gender)
```

```
[397]: df.head()
```

```
[397]:   Gender  Age  Annual Income (k$)  Spending Score (1-100)
0       1   19                15                39
1       1   21                15                81
2       0   20                16                 6
3       0   23                16               77
4       0   31                17               40
```

```
[398]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Gender                200 non-null   int64
1   Age                  200 non-null   int64
2   Annual Income (k$)    200 non-null   int64
3   Spending Score (1-100) 200 non-null   int64
dtypes: int64(4)
memory usage: 6.4 KB
```

4 Scaling Dataset values for uniformity

```
[399]: from sklearn.preprocessing import MinMaxScaler
```

```
[400]: df = pd.DataFrame(MinMaxScaler().fit_transform(df), columns=df.columns)
```


```
[401]: df.head()
```


```
[401]:   Gender  Age  Annual Income (k$)  Spending Score (1-100)
0     1.0  0.019231      0.000000      0.387755
1     1.0  0.057692      0.000000      0.816327
2     0.0  0.038462      0.008197      0.051020
3     0.0  0.096154      0.008197      0.775510
4     0.0  0.250000      0.016393      0.397959
```


5 Data Analysis, Outlier Detection & Outlier Elimination

```
[402]: sns.displot(df['Age'])  
sns.displot(df['Annual Income (k$)'])  
sns.displot(df['Spending Score (1-100)'])
```

```
[402]: <seaborn.axisgrid.FacetGrid at 0x7c71db8662f0>
```







```
[403]: sns.distplot(df['Age'])
```

<ipython-input-403-0fafa04ea3f6>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['Age'])
```

```
[403]: <Axes: xlabel='Age', ylabel='Density'>
```

```
[404]: sns.distplot(df['Annual Income (k$)'])
```

```
<ipython-input-404-5c9bfeb4bab1>:1: UserWarning:
```


```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```

```
Please adapt your code to use either `displot` (a figure-level function with  
similar flexibility) or `histplot` (an axes-level function for histograms).
```

```
For a guide to updating your code to use the new functions, please see  
https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751
```

```
sns.distplot(df['Annual Income (k$)'])
```

```
[404]: <Axes: xlabel='Annual Income (k$)', ylabel='Density'>
```



```
[405]: sns.distplot(df['Spending Score (1-100)'])
```

<ipython-input-405-beed7b40d5ab>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>


```
sns.distplot(df['Spending Score (1-100)'])
```

```
[405]: <Axes: xlabel='Spending Score (1-100)', ylabel='Density'>
```




```
[406]: sns.boxplot(df.Age)
```

```
[406]: <Axes: >
```




```
[407]: sns.boxplot(df['Spending Score (1-100)'])
```

```
[407]: <Axes: >
```



```
[408]: sns.boxplot(df['Annual Income (k$)'])
```

```
[408]: <Axes: >
```



```
[409]: Q1 = df['Annual Income (k$)'].quantile(0.25)
      Q3 = df['Annual Income (k$)'].quantile(0.75)
```


```
[410]: IQR = Q3 - Q1
      whisker_width = 1.5
```

```
[411]: lower_whisker = Q1 - (whisker_width*IQR)
      upper_whisker = Q3 + (whisker_width*IQR)
```

```
[412]: df['Annual Income (k$)'] = np.where(df['Annual Income (k$)'] > upper_whisker,
      ↪upper_whisker, np.where(df['Annual Income (k$)'] < lower_whisker,
      ↪lower_whisker, df['Annual Income (k$)']))
```


```
[413]: sns.boxplot(df['Annual Income (k$)'])
```

```
[413]: <Axes: >
```



```
[414]: plt.scatter(df['Annual Income (k$)'], df['Spending Score (1-100)'])
```

```
[414]: <matplotlib.collections.PathCollection at 0x7c71db350670>
```



```
[415]: X_train = df.drop(['Spending Score (1-100)'], axis=1)
       Y_train = df['Spending Score (1-100)']
```

```
[416]: X_train.head(), Y_train.head()
```

```
[416]: (   Gender      Age  Annual Income (k$)
0      1.0  0.019231          0.000000
1      1.0  0.057692          0.000000
2      0.0  0.038462          0.008197
3      0.0  0.096154          0.008197
4      0.0  0.250000          0.016393,
0      0.387755
1      0.816327
2      0.051020
3      0.775510
4      0.397959
Name: Spending Score (1-100), dtype: float64)
```


6 Finding Elbow Point (Possible 'K' value)

```
[417]: from sklearn.cluster import KMeans
```

```
[418]: k_rng = range(1,40)
sse = []

for k in k_rng:
    km = KMeans(n_clusters=k)
    km.fit(df[['Annual Income (k$)', 'Spending Score (1-100)']])
    sse.append(km.inertia_)
```

[illegible]


```

warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(

```

[419]: sse

```


[419]: [22.955815982449476,
13.908672064927337,
8.995430225375209,
6.055304436065896,
3.5287926379212995,
3.0667030475951083,
2.603378920756281,
2.187885090648535,
1.870681765951311,
1.65554733172543,
1.4601876173540953,
1.2648508391006437,
1.1345272740631709,
1.0315255123213856,
0.9552656259840362,
0.8832398005986595,
0.8364995074086296,
0.7597145770475486,
0.7125788697448221,

```

```
0.6661188281903597,  
0.6081924334679754,  
0.5868774306356352,  
0.5586492443579925,  
0.5255947055689602,  
0.4992976622795663,  
0.48431656019645386,  
0.45355573875679756,  
0.4349334492297876,  
0.4189017353108364,  
0.37528629789300777,  
0.369173284144015,  
0.3668613072634828,  
0.3457038348766507,  
0.31685030651587986,  
0.31137225211091013,  
0.287577555697374,  
0.27961820956283134,  
0.2697865474553408,  
0.2549457476255709]
```

```
[420]: plt.xlabel('K')  
plt.ylabel('Sum of Squared Error')  
plt.plot(k_rng, sse)
```

```
[420]: [<matplotlib.lines.Line2D at 0x7c71db3d0070>]
```



```
[421]: kmeans = KMeans(n_clusters=5)
kmeans
```

```
[421]: KMeans(n_clusters=5)
```

```
[422]: z = kmeans.fit_predict(df[['Annual Income (k$)', 'Spending Score (1-100)']])
z
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(
```

```
[422]: array([3, 4, 3, 4, 3, 4, 3, 4, 3, 4, 3, 4, 3, 4, 3, 4, 3, 4, 3, 4, 3, 4,
3, 4, 3, 4, 3, 4, 3, 4, 3, 4, 3, 4, 3, 4, 3, 4, 3, 0,
3, 4, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 1, 2, 0, 2, 1, 2, 1, 2,
0, 2, 1, 2, 1, 2, 1, 2, 1, 2, 0, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2,
1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2,
```

```
1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2,
1, 2], dtype=int32)
```

```
[423]: df_2 = df
```

```
[424]: df_2['Cluster'] = z
df_2.head()
```

```
[424]:
```

	Gender	Age	Annual Income (k\$)	Spending Score (1-100)	Cluster
0	1.0	0.019231	0.000000	0.387755	3
1	1.0	0.057692	0.000000	0.816327	4
2	0.0	0.038462	0.008197	0.051020	3
3	0.0	0.096154	0.008197	0.775510	4
4	0.0	0.250000	0.016393	0.397959	3

```
[425]: plt.figure(figsize=(10,8))


df1 = df_2[df_2.Cluster==0]
df2 = df_2[df_2.Cluster==1]
df3 = df_2[df_2.Cluster==2]
df4 = df_2[df_2.Cluster==3]
df5 = df_2[df_2.Cluster==4]

plt.scatter(df1['Annual Income (k$)'], df1['Spending Score (1-100)'],
            ↪label='Cluster 1')
plt.scatter(df2['Annual Income (k$)'], df2['Spending Score (1-100)'],
            ↪label='Cluster 2')
plt.scatter(df3['Annual Income (k$)'], df3['Spending Score (1-100)'],
            ↪label='Cluster 3')
plt.scatter(df4['Annual Income (k$)'], df4['Spending Score (1-100)'],
            ↪label='Cluster 4')
plt.scatter(df5['Annual Income (k$)'], df5['Spending Score (1-100)'],
            ↪label='Cluster 5')

plt.scatter(kmeans.cluster_centers_[ :,0], kmeans.cluster_centers_[ :,1],
            ↪color='black', marker='*', label='Centroid')

plt.xlabel('Annual Income (k$)')
plt.ylabel('Spending Score (1-100)')
plt.legend()
```

```
[425]: <matplotlib.legend.Legend at 0x7c71db2697e0>
```



7 Train - Test Split

```
[426]: df.drop(['Cluster'], axis=1, inplace=True)
```

```
[427]: df.head()
```

```
[427]:
```

	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1.0	0.019231	0.000000	0.387755
1	1.0	0.057692	0.000000	0.816327
2	0.0	0.038462	0.008197	0.051020
3	0.0	0.096154	0.008197	0.775510
4	0.0	0.250000	0.016393	0.397959

```
[428]: from sklearn.model_selection import train_test_split
```

```
[429]: X_train = df.drop(['Spending Score (1-100)'], axis=1)
       Y_train = df['Spending Score (1-100)']
```



```
[430]: # X_train, Y_train = train_test_split(df, test_size=0.5)
```

```
[431]: # print(X_train.shape)
# print(Y_train.shape)
```

```
[432]: # print(X_train.head())
```

8 Logistic Regression Modeling

```
[443]: from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
```

```
[444]: lr = LogisticRegression()
svc = SVC()
```

```
[435]: xtest, xtrain, ytest, ytrain = train_test_split(X_train, Y_train, test_size=0.
↳25)
```

```
[437]: xtest.shape, xtrain.shape
```

```
[437]: ((150, 3), (50, 3))
```

```
[441]: ytest.shape, ytrain.shape
```

```
[441]: ((150,), (50,))
```

I don't understand what to do next as I am not able to find the solution to the error

```
[447]: lr.fit(xtrain, ytrain)
```

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-447-dad6f18d5f54> in <cell line: 1>()
----> 1 lr.fit(xtrain, ytrain)

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py in
↳fit(self, X, y, sample_weight)
    1202         accept_large_sparse=solver not in ["liblinear", "sag",
↳"saga"],
    1203     )
-> 1204     check_classification_targets(y)
    1205     self.classes_ = np.unique(y)
    1206

/usr/local/lib/python3.10/dist-packages/sklearn/utils/multiclass.py in
↳check_classification_targets(y)
```

```
216         "multilabel-sequences",
217     ]:
--> 218         raise ValueError("Unknown label type: %r" % y_type)
219
220
```

```
ValueError: Unknown label type: 'continuous'
```

```
[ ]:
```