

# assignment-5

October 5, 2023

## 1 Kaggle Connection & DataFrame setup

```
[1049]: !pip install -q kaggle
```

```
[1050]: !mkdir ~/.kaggle
```

mkdir: cannot create directory '/root/.kaggle': File exists

```
[1051]: !cp kaggle.json ~/.kaggle
```

cp: cannot stat 'kaggle.json': No such file or directory

```
[1052]: ! kaggle datasets download -d vjchoudhary7/  
↪customer-segmentation-tutorial-in-python
```

Warning: Your Kaggle API key is readable by other users on this system! To fix this, you can run 'chmod 600 /root/.kaggle/kaggle.json'  
customer-segmentation-tutorial-in-python.zip: Skipping, found more recently modified local copy (use --force to force download)

```
[1053]: !unzip /content/customer-segmentation-tutorial-in-python.zip
```

Archive: /content/customer-segmentation-tutorial-in-python.zip  
inflating: Mall\_Customers.csv

## 2 Pre-Processing

```
[1054]: import pandas as pd  
import numpy as np  
import seaborn as sns  
import matplotlib.pyplot as plt
```

```
[1055]: df = pd.read_csv('./Mall_Customers.csv')  
df.head()
```

```
[1055]:
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	39

1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40

```
[1056]: df.describe()
```

```
[1056]:
```

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)
count	200.000000	200.000000	200.000000	200.000000
mean	100.500000	38.850000	60.560000	50.200000
std	57.879185	13.969007	26.264721	25.823522
min	1.000000	18.000000	15.000000	1.000000
25%	50.750000	28.750000	41.500000	34.750000
50%	100.500000	36.000000	61.500000	50.000000
75%	150.250000	49.000000	78.000000	73.000000
max	200.000000	70.000000	137.000000	99.000000

```
[1057]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CustomerID            200 non-null   int64
1   Gender                200 non-null   object
2   Age                   200 non-null   int64
3   Annual Income (k$)    200 non-null   int64
4   Spending Score (1-100) 200 non-null   int64
dtypes: int64(4), object(1)
memory usage: 7.9+ KB
```

```
[1058]: df.isnull().values.any()
```

```
[1058]: False
```

```
[1059]: df.shape
```

```
[1059]: (200, 5)
```

```
[1060]: # Dropping 'CustomerID' as it has no impact or connection to dataset or data
        ↪ values
df.drop(['CustomerID'], axis=1, inplace=True)
```


```
[1061]: sns.heatmap(df.corr(), annot=True)
```

```
<ipython-input-1061-6dc1c4c1753e>:1: FutureWarning: The default value of
numeric_only in DataFrame.corr is deprecated. In a future version, it will
```

default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.


```
sns.heatmap(df.corr(), annot=True)
```

[1061]: <Axes: >



```
[1062]: sns.pairplot(df)
```

[1062]: <seaborn.axisgrid.PairGrid at 0x7dadda28db40>



### 3 Converting Categorical Data (Columns) to Numerical

```
[1063]: df['Gender'].value_counts()
```

```
[1063]: Female    112
        Male      88
        Name: Gender, dtype: int64
```

```
[1064]: from sklearn.preprocessing import LabelEncoder
        le = LabelEncoder()
```

```
[1065]: # Label Encoding 'Gender' column
# '1' == 'Male' && 0 == 'Female'
df['Gender'] = le.fit_transform(df.Gender)
```

```
[1066]: df.head()
```

```
[1066]:
```

	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	19	15	39
1	1	21	15	81
2	0	20	16	6
3	0	23	16	77
4	0	31	17	40


```
[1067]: df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Gender                                200 non-null   int64
1   Age                                   200 non-null   int64
2   Annual Income (k$)                   200 non-null   int64
3   Spending Score (1-100)               200 non-null   int64
dtypes: int64(4)
memory usage: 6.4 KB
```


## 4 Data Analysis, Outlier Detection & Outlier Elimination

```
[1068]: sns.displot(df['Age'])
sns.displot(df['Annual Income (k$)'])
sns.displot(df['Spending Score (1-100)'])
```

```
[1068]: <seaborn.axisgrid.FacetGrid at 0x7dadd99c8370>
```







```
[1069]: sns.distplot(df['Age'])
```

<ipython-input-1069-0fafa04ea3f6>:1: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.


Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['Age'])
```

```
[1069]: <Axes: xlabel='Age', ylabel='Density'>
```





```
[1070]: sns.distplot(df['Annual Income (k$)'])
```

<ipython-input-1070-5c9bfeb4bab1>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['Annual Income (k$)'])
```

```
[1070]: <Axes: xlabel='Annual Income (k$)', ylabel='Density'>
```



```
[1071]: sns.distplot(df['Spending Score (1-100)'])
```

<ipython-input-1071-beed7b40d5ab>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>


```
sns.distplot(df['Spending Score (1-100)'])
```

```
[1071]: <Axes: xlabel='Spending Score (1-100)', ylabel='Density'>
```



```
[1072]: sns.boxplot(df['Annual Income (k$)'])
```

```
[1072]: <Axes: >
```



```
[1073]: Q1 = df['Annual Income (k$)'].quantile(0.25)
        Q3 = df['Annual Income (k$)'].quantile(0.75)
```


```
[1074]: IQR = Q3 - Q1
        whisker_width = 1.5
```

```
[1075]: lower_whisker = Q1 - (whisker_width*IQR)
        upper_whisker = Q3 + (whisker_width*IQR)
```

```
[1076]: df['Annual Income (k$)'] = np.where(df['Annual Income (k$)'] > upper_whisker,
        ↪upper_whisker, np.where(df['Annual Income (k$)'] < lower_whisker,
        ↪lower_whisker, df['Annual Income (k$)']))
```


```
[1077]: sns.boxplot(df['Annual Income (k$)'])
```

```
[1077]: <Axes: >
```



```
[1078]: plt.scatter(df['Annual Income (k$)'], df['Spending Score (1-100)'])
```

```
[1078]: <matplotlib.collections.PathCollection at 0x7dadd943d780>
```



## 5 Finding Elbow Point (Possible 'K' value)

```
[1079]: from sklearn.cluster import KMeans
```

```
[ ]: k_rng = range(1,40)
sse = []

for k in k_rng:
    km = KMeans(n_clusters=k)
    km.fit(df[['Annual Income (k$)', 'Spending Score (1-100)']])
    sse.append(km.inertia_)
```

```
[ ]: sse
```

```
[1082]: plt.xlabel('K')
plt.ylabel('Sum of Squared Error')
plt.plot(k_rng, sse)
```

```
[1082]: [<matplotlib.lines.Line2D at 0x7dadd94d1300>]
```



## 6 K-Means

kmeans

```
KMeans(n_clusters=5)
```

Z

```

/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(

```

```
array([4, 3, 4, 3, 4, 3, 4, 3, 4, 3, 4, 3, 4, 3, 4, 3, 4, 3, 4, 3,
       4, 3, 4, 3, 4, 3, 4, 3, 4, 3, 4, 3, 4, 3, 4, 3, 4, 2,
       4, 3, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
       2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
       2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
       2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 1, 0, 1, 2, 1, 0, 1, 0, 1])
```

```

2, 1, 0, 1, 0, 1, 0, 1, 0, 1, 2, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1,
0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1,
0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1,
0, 1], dtype=int32)

```

```
[1085]: df_2 = df
```

```
[1086]: df_2['Cluster'] = z
df_2.head()
```

```
[1086]:
```

	Gender	Age	Annual Income (k\$)	Spending Score (1-100)	Cluster
0	1	19	15.0	39	4
1	1	21	15.0	81	3
2	0	20	16.0	6	4
3	0	23	16.0	77	3
4	0	31	17.0	40	4

```
[1087]: plt.figure(figsize=(10,8))

df1 = df_2[df_2.Cluster==0]
df2 = df_2[df_2.Cluster==1]
df3 = df_2[df_2.Cluster==2]
df4 = df_2[df_2.Cluster==3]
df5 = df_2[df_2.Cluster==4]


plt.scatter(df1['Annual Income (k$)'], df1['Spending Score (1-100)'],
            label='Cluster 1')
plt.scatter(df2['Annual Income (k$)'], df2['Spending Score (1-100)'],
            label='Cluster 2')
plt.scatter(df3['Annual Income (k$)'], df3['Spending Score (1-100)'],
            label='Cluster 3')
plt.scatter(df4['Annual Income (k$)'], df4['Spending Score (1-100)'],
            label='Cluster 4')
plt.scatter(df5['Annual Income (k$)'], df5['Spending Score (1-100)'],
            label='Cluster 5')

plt.scatter(kmeans.cluster_centers_[ :,0], kmeans.cluster_centers_[ :,1],
            color='black', marker='*', label='Centroid')

plt.xlabel('Annual Income (k$)')
plt.ylabel('Spending Score (1-100)')
plt.legend()
```

```
[1087]: <matplotlib.legend.Legend at 0x7dadd9366620>
```





## 7 Scaling & Train - Test Split

```
[1088]: from sklearn.cluster import KMeans
        from sklearn.preprocessing import MinMaxScaler
        from sklearn.model_selection import train_test_split
```

```
[1089]: df.drop(['Cluster'], axis=1, inplace=True)
```

```
[1090]: df = pd.DataFrame(MinMaxScaler().fit_transform(df), columns=df.columns)
```

```
[1091]: # X = df.drop('Spending Score (1-100)', axis=1)
        # y = df['Spending Score (1-100)']
```

```
[1092]: # xtrain, xtest, ytrain, ytest = train_test_split(X, y, test_size=0.3)
```

```
[1098]: kmeans = KMeans(n_clusters=5, init = 'k-means++', random_state=0)
```

```
[1099]: kmeans.fit(df)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(
```

```
[1099]: KMeans(n_clusters=5, random_state=0)
```

```
[1100]: X.head()
```

```
[1100]:
```

	Gender	Age	Annual Income (k\$)
0	1.0	0.019231	0.000000
1	1.0	0.057692	0.000000
2	0.0	0.038462	0.008493
3	0.0	0.096154	0.008493
4	0.0	0.250000	0.016985

```
[1101]: kmeans.predict(df)
```

```
[1101]: array([3, 3, 2, 1, 1, 1, 2, 1, 4, 1, 4, 1, 2, 1, 4, 3, 2, 3, 4, 1, 4, 3,
        2, 3, 2, 3, 2, 3, 2, 1, 4, 1, 4, 3, 2, 1, 2, 1, 2, 1, 2, 3, 4, 1,
        2, 1, 2, 1, 1, 1, 2, 3, 1, 4, 2, 4, 2, 4, 1, 4, 4, 3, 2, 2, 4, 3,
        2, 2, 3, 1, 4, 2, 2, 2, 4, 3, 2, 4, 1, 2, 4, 3, 4, 2, 1, 4, 2, 1,
        1, 2, 2, 3, 4, 2, 1, 3, 2, 1, 4, 3, 1, 2, 4, 3, 4, 1, 2, 4, 4, 4,
        4, 1, 2, 3, 1, 1, 2, 2, 2, 2, 3, 2, 1, 3, 1, 1, 0, 3, 4, 3, 0, 3,
        1, 1, 0, 1, 2, 3, 0, 1, 2, 3, 1, 1, 0, 3, 4, 1, 2, 3, 0, 3, 2, 1,
        2, 1, 0, 1, 0, 1, 2, 1, 0, 1, 0, 1, 0, 1, 2, 3, 0, 3, 0, 3, 2, 1,
        0, 3, 0, 3, 2, 1, 0, 1, 2, 3, 2, 3, 2, 1, 2, 1, 0, 1, 2, 1, 2, 3,
        0, 3], dtype=int32)
```

## 8 Prediction

```
[1102]: kmeans.predict([[1, 19, 15, 81]])[0]
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does
not have valid feature names, but KMeans was fitted with feature names
warnings.warn(
```

```
[1102]: 3
```